

Effect of Replacing Attention Heads on Downstream Model Evaluations (Chosen Program vs. Wrong-Category baseline, mean $\pm$ std across seeds)

